

ADARI DURGA SAI SHASHANK

Lead Data Engineer | Azure Lakehouse & Real-Time Streaming | AI-Augmented Engineering: LLM Agents, MCP, AI Guardrails

Siddipet, Telangana, India | +91 9542955836 | saishashankadari@gmail.com | [linkedin.com/in/adari-shashank](https://www.linkedin.com/in/adari-shashank) | github.com/adarishashank

Professional Summary

Lead Data Engineer with 5+ years building batch and real-time data platforms on Microsoft Azure: Databricks, Delta Lake, Unity Catalog, Microsoft Fabric, Data Factory, Event Hubs and Kusto. Leads 3 concurrent PepsiCo migrations from Hive Metastore and Teradata to Databricks, including a sync framework designed to land a 51-billion-row table within a 1-hour window. Early AI adopter who evaluates frontier LLMs and rolls out guardrailed AI-assisted engineering to project teams.

Technical Skills

- **Languages:** Python, PySpark, SQL (Spark SQL, T-SQL), KQL, PowerShell, JavaScript
- **Lakehouse & Platforms:** Azure Databricks, Delta Lake, Unity Catalog, Microsoft Fabric, Power BI, ADLS Gen2, Synapse Analytics, Azure Data Explorer (Kusto), Teradata
- **Streaming:** Spark Structured Streaming, Apache Kafka, Event Hubs, Stream Analytics, IoT Hub, Event Grid
- **Orchestration & DevOps:** Azure Data Factory, Databricks Workflows, Azure DevOps (YAML pipelines), ARM templates, Databricks Asset Bundles, Git
- **Data Engineering:** medallion architecture, dimensional modeling, incremental MERGE loads, row-hash reconciliation, Unity Catalog governance, Teradata and Hive Metastore migration
- **Observability:** Azure Monitor, Log Analytics, Databricks system tables, Grafana, Datadog
- **AI & LLM:** LLM evaluation (Claude, ChatGPT, Grok, Mistral; System 1 and System 2 models), LLM agents, Model Context Protocol (MCP), agentic workflows, prompt engineering, AI guardrails

Professional Experience

Technical Lead, Data Engineering | Cliff Systems LLP | Hyderabad, India | Aug 2026 – Present

Client: PepsiCo

- Lead 3 concurrent migration programs: Hive Metastore to Unity Catalog for a pan-European commercial analytics platform and a South African topline / COGS pipeline, plus an APAC Teradata-to-Databricks wave; own architecture, standards, CI/CD and production sign-off
- Architected a config-driven Unity Catalog-to-ADLS sync framework (ADF orchestrator, JSON registry, MERGE and full-refresh modes, run-overlap protection) designed to land a 51-billion-row fact table within a 1-hour window
- Built the production sign-off framework: schema diffs, row counts and SHA-256 row-hash reconciliation across 2 workspaces with a 5% variance gate, producing audit-ready evidence for the independent validation team
- Re-engineered ARM template parameterization from 739 parameters to within Azure's 256-parameter limit, unblocking ADF CI/CD, while leading the split of 1 factory into 2 regional factories (UK and US)
- Defined the migration playbook (external-location governance, clone-versus-CTAS strategy, Delta log recovery, T-SQL-to-Spark parity rules) and authored a 6-volume Databricks enablement series
- Drove AI adoption across project teams: evaluated 4+ frontier LLM families (Claude, ChatGPT, Grok, Mistral) across System 1 and System 2 reasoning models, then rolled out AI-assisted coding, review and audits with guardrails such as restricted Databricks data access and human review

Senior Analyst, Full-Stack Engineering | Accenture | Hyderabad, India | Aug 2021 – Jul 2026

Clients: Mars Global and a Fortune 500 global manufacturer; progressed from platform support to development to migration engineering

Client: Fortune 500 Global Manufacturer — Teradata to Azure Databricks Migration | Mar 2025 – Jul 2026

- Migrated Teradata workloads across 4 regions (Europe, APAC, Africa / Middle East / South Asia, North America) to Databricks and Unity Catalog, converting BTEQ and KSH jobs into PySpark and Spark SQL medallion pipelines, including MERGE loads for SAP BW and JDA / Blue Yonder feeds
- Orchestrated dependencies across 23 Databricks workflows with event-driven ADF and Event Grid gates that verify upstream loads via the ADF REST API, replaced KSH and BTEQ schedulers with Delta control tables, and hardened CI/CD with PowerShell retry logic
- Owned source-to-target parity through go-live, resolving Teradata-to-Spark semantic gaps (case sensitivity, numeric overflow) and Delta write-concurrency conflicts with documented root-cause analysis

Client: Mars Global — Digital Twin for Manufacturing (Real-Time IoT Platform) | Sep 2022 – Feb 2025

- Delivered real-time IoT telemetry pipelines (IoT Hub, Event Hubs, Stream Analytics) into Azure Data Explorer and SQL, plus ADF batch ETL and Logic Apps alerting, for a Digital Twin platform that cut material wastage by 35%
- Spearheaded an Azure PaaS optimization that reduced cost by 20% and improved performance by 30%, automating deployments with Python and monitoring with Grafana and Datadog

Client: Mars Global — Azure PaaS Engineering and Support | Aug 2021 – Aug 2022

- Resolved 700 incidents and service requests within SLA for an enterprise Azure PaaS estate, documenting root causes to prevent recurrence

Education

Bachelor of Technology, Computer Science & Engineering — Koneru Lakshmaiah Education Foundation (KL University), Guntur, India | 2017 – 2021

Certifications

- Microsoft Certified: Azure Administrator Associate (AZ-104) — 2023 – 2025

Awards & Leadership

- **University Innovation Fellow**, Stanford d.school (2020 cohort) — introduced a Week of Welcome for first-year students, a peer-learning program and campus process improvements
- **Accelerated Promotion Award, Accenture** — recognized by the global management committee; also named Buddy Champion for mentoring new joiners